Representation Theory and Locality in Celestial Holography

Mathew Calkins

May 1, 2026

Abstract

An introduction to celestial holography, leading up to some recent research on multi-collinear singular structures.

Contents

1 Conventions **1**

2 Motivation **1**
2.1 AdS/CFT . . . 1
2.2 Asymptotically Flat Space . . . 1
2.3 Troubles with de Sitter . . . 1

3 Kinematics of The Celestial Dual **2**
3.1 Canonical and Non-Canonical Group Isomorphisms . . . 2
3.2 Mellin-Transformation Diagonalizes Boosts . . . 5
3.3 Massive Particles . . . 6

4 Dynamics of the Celestial Dual **6**
4.1 OPEs Organized by Symmetry . . . 6
4.2 OPEs from Bulk Scattering . . . 7

1 Conventions

In Minkowski space we use the mostly-plus metric

$$\eta = \mathrm{diag}(-1, +1, \cdots, +1). \tag{1}$$

2 Motivation

TODO

2.1 AdS/CFT

TODO

2.2 Asymptotically Flat Space

TODO

2.3 Troubles with de Sitter

TODO

3 Kinematics of The Celestial Dual

Though we have discussed generally asymptotically flat spacetimes, we now restrict our attention to flat Minkowski space. We recall the well-known group isomorphism between the Lorentz group in signature $(d+1,1)$ and the Euclidean conformal group in d dimensions. At $d=2$ this becomes

$$SO(3,1) \simeq \mathrm{Conf}(2,0). \tag{2}$$

Here and elsewhere, when we refer to a Lie group we will always mean the its connected component containing the identity. Properly speaking we are therefore referring to the proper orthochronous Lorentz group $SO^{\uparrow}(3,1)$. On the other side of our isomorphism we are excluding orientation-reversing transformations; using the standard complex-analytic language we are equivalently restricting to holomorphic functions. Also as usual, the Euclidean conformal group in d dimensions should be thought of not as a group of automorphisms of real Euclidean space $\mathbb{R}^d$, but of its conformal compactification S^d. TODO: Implement this on the Lie-algebra level?

Inspired by the symmetry-matching game we played earlier for AdS/CFT, taking this isomorphism seriously leads one to wonder whether physics in Minkowski space might have a dual 2D conformal field theory. If we want this to lift to more general asymptotically flat spacetimes and have a shot and describing bulk quantum gravity, the dictionary ought to be built from bulk ingredients that make sense in such a theory.

The simplest candidate is the S-matrix, and so far this appears to be exactly the right object to use. An n-point amplitude has functional dependence

$$\mathcal{A} \sim \mathcal{A}_n\left(p_1, \sigma_1, n_1; \ldots; p_n, \sigma_n, n_n\right) \tag{3}$$

where, following the notation of [**?**], the ith particle has momentum p_i, spin/helicity labeled by σ_i, and flavor indices given by n_i. The idea behind the celestial dictionary is to parametrize our momenta and apply an appropriate change of basis after which the RHS to produce an object which transforms under Lorentz transformations (equivalently, under global conformal transformations) like the n-point vacuum correlator of a list of global conformal primaries. That is, we want

$$\mathcal{A}_n\left(p_1, \sigma_1, n_1; \ldots; p_n, \sigma_n, n_n\right) \xmapsto{\text{change of basis?}} \left\langle \mathcal{O}_{h_1,\bar{h}_1}(z_1,\bar{z}_1) \cdots \mathcal{O}_{h_n,\bar{h}_n}(z_n,\bar{z}_n)\right\rangle \tag{4}$$

where the RHS is to be computed in some dual theory on the celestial sphere S^2. Now the whole game of scattering is that amplitudes can be computed as matrix elements between a set of m incoming and k outgoing states for $m+k=n$.

$$\mathcal{A}_n\left(p_1, \sigma_1, n_1; \ldots; p_n, \sigma_n, n_n\right) = \left\langle p_1^{\text{in}}, \sigma_1^{\text{in}}, n_1^{\text{in}}; \ldots; p_m^{\text{in}}, \sigma_m^{\text{in}}, n_m^{\text{in}}; \middle| S \middle| p_1^{\text{out}}, \sigma_1^{\text{out}}, n_1^{\text{out}}; \ldots; p_k^{\text{out}}, \sigma_k^{\text{out}}, n_k^{\text{out}}\right\rangle \tag{5}$$

By construction these transform under Lorentz as the tensor product of single-particle states, so it suffices to find an appropriate change of basis on single-particle basis states

$$|p,\sigma\rangle \xmapsto{\text{change of basis?}} \left|z_1, \bar{z}_1, h_1, \bar{h}_1\right\rangle \tag{6}$$

Before we work out the details of this change of basis, we make a slightly embarrassing admission: at the time of writing, it is not yet clear how to make this story work correctly when any of our bulk particles are massive. We can play the same representation-theoretic game, but the end result turns out not have the locality properties we would expect from a Euclidean n-point correlator. We will describe this problem more in detail later, but for now we restrict to the case that all off-shell momenta are massless.

3.1 Canonical and Non-Canonical Group Isomorphisms

In order to construct the celestial dictionary, we begin by spelling out the relevant group isomorphisms.

$$\begin{tikzcd}
SO^{+}(3,1) \arrow[r, leftrightsquigarrow] \arrow[d, leftrightarrow] & SL(2,\mathbb{C})/\mathbb{Z}_2 \arrow[d, leftrightarrow] \\
\mathrm{Conf}^{+}(S^2) \arrow[r, leftrightsquigarrow] & \mathrm{Conf}^{+}(\mathbb{C}\cup\{\infty\})
\end{tikzcd}$$

We must first distinguish between the conformal automorphisms of the sphere from those the compactified plane. We write these groups respectively as $\mathrm{Conf}(S^2)$ and $\mathrm{Conf}(\mathbb{C} \cup \{\infty\})$, and use a superscript $^+$ to denote the connected component of any Lie group containing the identity (which for these groups is exactly the subgroup of orientation-preserving such maps). Of course the manifolds S^2 and $\mathbb{C} \cup \{\infty\}$ are conformally equivalent, but not canonically so. Equivalently, there are many many to endow S^2 with conformal coordinates $(z, \bar{z})$ so that the metric on S^2 is some local rescaling of $dz \wedge d\bar{z}$. In the commutative diagram above, we write a squiggly bidirectional arrow to indicate that $\mathrm{Conf}(S^2)$ and $\mathrm{Conf}(\mathbb{C} \cup \infty\}$ are non-canonically isomorphic groups. More quantitatively, we can describe the degree of non-canonicality by noting that any two conformal maps

$$\phi, \phi' : \mathbb{C} \cup \{\infty\} \to S^2 \tag{7}$$

differ exactly by composition with a conformal transformation of S^2.

$$\phi' = \psi \circ \phi, \quad \psi = \phi' \circ \phi^{-1} : S^2 \to S^2. \tag{8}$$

So the ambiguity is exactly the group of canonical automorphisms of S^2, which form a Lie group of real dimension 6.

Now for the top-most edge of our diagram. We recall that $SO^+(3,1)$ has a universal cover of degree 2 isomorphic to $SL(2,\mathbb{C})$, or equivalently that $SO^+(3,1) \simeq SL(2,\mathbb{C})/\mathbb{Z}_2$. The usual construction works by introducing the Pauli matrices

$$\begin{aligned} \left(\sigma^0_{\alpha\dot{\alpha}}\right) &= \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}, \quad \left(\sigma^1_{\alpha\dot{\alpha}}\right) = \begin{bmatrix} 0 & 1 \\ 1 & 0 \end{bmatrix}, \\ \left(\sigma^2_{\alpha\dot{\alpha}}\right) &= \begin{bmatrix} 0 & -i \\ i & 0 \end{bmatrix}, \quad \left(\sigma^3_{\alpha\dot{\alpha}}\right) = \begin{bmatrix} 1 & 0 \\ 0 & -1 \end{bmatrix} \end{aligned} \tag{9}$$

then building the natural mapping

$$\rho : \mathbb{R}^{3+1} \to \mathrm{Her}(2,\mathbb{C}) := \{\text{Hermitian } 2 \times 2 \text{ matrices}\} \tag{10}$$

given by

$$\rho(p)_{\alpha\dot{\alpha}} = \sigma^\mu_{\alpha\dot{\alpha}} \eta_{\mu\nu} p^\nu \equiv p_{\alpha\dot{\alpha}} \tag{11}$$

or in components

$$\left(p_{\alpha\dot{\alpha}}\right) = \begin{bmatrix} -p^0 + p^3 & p^1 - ip^2 \\ p^1 + ip^2 & -p^0 - p^3 \end{bmatrix}. \tag{12}$$

Now by direct computation we see that this obeys

$$\det\left(p_{\alpha\dot{\alpha}}\right) = \left(p^0\right)^2 - \left(p^1\right)^2 - \left(p^2\right)^2 - \left(p^3\right)^2 = -p \cdot p. \tag{13}$$

So the proper orthochronous Lorentz transformations on p correspond to the determinant- and Hermiticity-preserving linear maps on $\rho(p)$. The latter are exactly the conjugations

$$M \mapsto AMA^\dagger \text{ for } A \in SL(2,\mathbb{C}) \tag{14}$$

so any $A \in SL(2,\mathbb{C})$ has an associated proper orthchronous Lorentz transformation $\Lambda \in SO^+(3,1)$ defined implicitly by the rule

$$A\rho(p)A^\dagger = \rho(\Lambda p) \text{ for all } p \in \mathbb{R}^{3+1}. \tag{15}$$

This mapping $A \mapsto \Lambda$ is exactly the projection map establishing $SL(2,\mathbb{C})$ as the double cover of $SO^+(3,1)$. But we made a non-canonical choice: we chose (σ_μ) as a real basis for the set of all Hermitian 2×2 matrices satisfying the determinant condition (13). But this choice is non-unique; with a bit of thought, one can see that some other basis $(\hat{e}_\mu)$ has this property iff

$$\hat{e}^\mu = \Lambda^\mu{}_\nu \hat{e}^\nu \tag{16}$$

for Λ a Lorentz transformation. Again, we see that the degree of ambiguity is parametrized by a Lie group of real dimension 6.

On the other hand, the isomorphisms indicated by solid bidirectional arrows are canonical. Along the rightmost edge of our diagram, we recall that the the conformal automorphisms of $\mathbb{C} \cup \{\infty\}$ connected to the identity are exactly the fractional linear transformations

$$z \mapsto \frac{az+b}{cz+d} \tag{17}$$

for $a, b, c, d \in \mathbb{C}$ satisfying WLOG $ad - bc = 1$. So to any special linear matrix we can associate such a mapping as

$$\begin{bmatrix} a & b \\ c & d \end{bmatrix} \mapsto \left(z \mapsto \frac{az+b}{cz+d} \right). \tag{18}$$

This group homomorphism is 2:1, with $\pm M$ having the same image, hence induces a canonical isomorphism between $\mathrm{Conf}^+(\mathbb{C} \cup \{\infty\})$ and $SL(2, \mathbb{C})/\mathbb{Z}_2$.

The canonical isomorphism between $SO^+(3,1)$ and $\mathrm{Conf}(S^2)$ is best described in more geometrical terms. We associate to any nonzero null vector $p = (p^0, \vec{p}) \in \mathbb{R}^{3+1}$ the normalized future-directed null vector $\hat{p}$ by taking

$$p = p^0 \hat{p} \implies \hat{p} = \left(1, \frac{\vec{p}}{p^0}\right) \text{ where } \frac{\vec{p}}{p^0} \in S^2 \tag{19}$$

The unit vectors $\hat{p}$ correspond naturally to points on S^2, and they get moved around if we replace p with Λp. So any Lorentz transformation $\Lambda \in SO(3,1)$ induces an automorphism (which one can convince oneself is conformal) of S^2 given by

$$\hat{p} \mapsto \widehat{\Lambda p}. \tag{20}$$

This map is well-defined, as it doesn't depend on how at intermediate steps we lift $\hat{p}$ to p. It also illustrates exactly why the correct isomorphism is between the proper orthochronous Lorentz Lorentz transformations $\Lambda \in SO^+(3,1)$ and orientation-preserving conformal transformations $f \in \mathrm{Conf}^+(S^2)$: any Lorentz transformation which isn't proper orthochronous must be a proper orthochronous one composed with the PT reflection $x^\mu \mapsto -x^\mu$, but such a transformation corresponds to the identity transformation on S^2. We see similarly that any Lorentz transformation built from only one of P, T is orientation-reversing on S^2. So we can lift to a slightly more detailed homomorphism from Minkowski-norm-preserving linear transformations to not-necessarily-orientation-preserving conformal maps

$$\pi : O(3,1) \to \mathrm{Conf}(S^2). \tag{21}$$

The domain has four connected components corresponding to whether a transformation includes P and T, the codomain has two corresponding to whether a transformation flips orientation, and the map is 2:1 with kernel PT.

Now we want to parametrize the null external momenta appearing in our scattering amplitude in a way which plays nicely with these isomorphisms. Concretely, we want to put conformal coordinates $(z, \bar{z})$ on $S^2 \simeq \{\text{normalized null momenta}\}$. For reasons just described this will require some arbitrary choices, but we know that z should transform under a fractional-linear mapping when p transforms under a linear one, and we know that any two momenta differing by overall scaling should correspond to the same point on S^2. We therefore follow [**?**] and take

$$z = \frac{p^1 + ip^2}{p^0 + p^3} \tag{22}$$

From this we see immediately that a spatial rotation about the bulk x^3-axis and boosts along the same corresponding on the boundary to rotation and dilatations centered at the origin. Equivalently, this corresponds to putting coordinates on S^2 via a stereographic projection onto the x^1x^2-plane.

After some algebraic fiddling modulo the relation $p_\mu p^\mu = 0$ we find

$$\left(\frac{p^1}{p^0}, \frac{p^2}{p^0}, \frac{p^3}{p^0}\right) = \left(\frac{z+\bar{z}}{1+z\bar{z}}, \frac{-i(z-\bar{z})}{1+z\bar{z}}, \frac{1-z\bar{z}}{1+z\bar{z}}\right). \tag{23}$$

So up to an arbitrary choice of normalization (which we make slightly differently than we did above), we decompose

$$p(\omega, z, \bar{z}) = \omega \hat{p}(z, \bar{z}) = \omega \left(1 + z\bar{z}, z + \bar{z}, -i(z - \bar{z}), 1 - z\bar{z}\right). \tag{24}$$

3.2 Mellin-Transformation Diagonalizes Boosts

Now given a single-particle state of momentum $p = \omega\hat{p}(z,\bar{z})$ and helicity H, we can write

$$|p, H\rangle \equiv |\omega, z, \bar{z}, H\rangle \,. \tag{25}$$

We noted above that our goal is to produce from this ket another that transforms under global conformal transformations "like a global conformal primary." What exactly does this mean? Algebraically, we know [**?**] that under a global transformation

$$z \mapsto z' = \frac{az+b}{cz+d}, \quad ad - bc = 1 \tag{26}$$

a global conformal primary in a 2D CFT with conformal weights $(h, \bar{h})$ transforms as

$$\mathcal{O}'_{h,\bar{h}}(z', \bar{z}') = \left(\frac{dz'}{dz}\right)^{-h} \left(\frac{d\bar{z}'}{d\bar{z}}\right)^{-\bar{h}} \mathcal{O}_{h,\bar{h}}(z,\bar{z}) = (cz+d)^{-2h}(\bar{c}\bar{z}+\bar{d})^{-2\bar{h}}\mathcal{O}_{h,\bar{h}}(z,\bar{z}). \tag{27}$$

In terms of Lie algebra elements, $\mathcal{O}_{h,\bar{h}}(z,\bar{z})$ simultaneously diagonalizes the subgroup of $\mathrm{Conf}(\mathbb{C} \cup \{\infty\})$ which leaves $z, \bar{z}$ fixed, *i.e.*, the subgroup generated by rotations and dilatations centered at $(z,\bar{z})$. Inserting our operators at the origin for ease of notation (the general case can be recovered by conjugation by translation), under a rotation $z' = e^{i\theta}z$ we have $(a,b,c,d) = (e^{i\theta/2}, 0, 1, e^{-i\theta/2})$ hence

$$\mathcal{O}'_{h,\bar{h}}(0,0) = e^{-i\theta(h-\bar{h})}\mathcal{O}_{h,\bar{h}}(0,0) \equiv e^{-i\theta s}\mathcal{O}_{h,\bar{h}}(0,0). \tag{28}$$

while under under an origin-centered dilatation $z' = e^{\Omega}z$ we have

$$\mathcal{O}'_{h,\bar{h}}(0,0) = e^{-\Omega(h+\bar{h})}\mathcal{O}'_{h,\bar{h}}(0,0) \equiv e^{-\Omega\Delta}\mathcal{O}'_{h,\bar{h}}(0,0) \tag{29}$$

So as usual we identify the combinations

$$(\Delta, s) = \left(h+\bar{h}, h-\bar{h}\right) \tag{30}$$

as the conformal scaling weight and spin. By conjugation under translation, we see that in general $\mathcal{O}_{h,\bar{h}}(z,\bar{z})$ simultaneously diagonalized rotations around and dilatations about $(z,\bar{z})$. To match this in bulk Lorentz language, we need a $p(z,\bar{z})$-dependent ket which simultaneously diagonalizes rotation about and boosts along $p(z,\bar{z})$. We have already accomplished half of this: under a spatial rotation by θ about the spatial part of p, precisely by the definition of helicity we have

$$|\omega, z, \bar{z}, H\rangle \xmapsto{R_z(\theta)} e^{-i\theta H}\,|\omega, z, \bar{z}, H\rangle \,. \tag{31}$$

So we have already diagonalized boundary rotations centered at $(z,\bar{z})$ and identify the bulk helicity H with the boundary conformal spin s. We have not yet diagonalized boosts along p, under which our ket transforms as

$$|\omega, z, \bar{z}, H\rangle \xmapsto{B_z(\beta)} \left|e^{\beta}\omega, z, \bar{z}, H\right\rangle \tag{32}$$

where β is the rapidity of the boost. We can reverse-engineer what integral transform does the job: we define a new ket related to the old one by a Mellin-transform with respect to the scaling parameter ω.

$$|\Delta, z, \bar{z}, s\rangle = \int_0^\infty \frac{d\omega}{\omega}\omega^{\Delta}\,|\omega, z, \bar{z}, H = s\rangle \,. \tag{33}$$

Then under a boost we have

$$\begin{aligned}
|\Delta, z, \bar{z}, s\rangle &\xmapsto{B_z(\beta)} \int_0^\infty \frac{d\omega}{\omega}\omega^{\Delta}\left|e^{\beta}\omega, z, \bar{z}, H = s\right\rangle \\
&= \int_0^\infty \frac{d\left(e^{-\beta}\omega\right)}{\left(e^{-\beta}\omega\right)}\left(e^{-\beta}\omega\right)^{\Delta}\left|e^{\beta}\left(e^{-\beta}\omega\right), z, \bar{z}, H = s\right\rangle \\
&= e^{-\beta\Delta}\int_0^\infty \frac{d\omega}{\omega}\omega^{\Delta}\,|\omega, z, \bar{z}, H = s\rangle \\
&= e^{-\beta\Delta}\,|\Delta, z, \bar{z}, s\rangle \,.
\end{aligned} \tag{34}$$

TODO: Describe why these satisfy the right highest-weight condition to be conformal primaries. So far our argument is consistent with them being descendants.

Finally, we can write down our proposed celestial dictionary. Given a bulk scattering amplitude between n on-shell external particles each with null momentum $p = \omega\hat{p}(z,\bar{z})$, we integrate out each ω by a Mellin transform to product an object which, at least kinematically, behaves like an n-point correlator in a 2D CFT.

$$\langle \mathcal{O}_{h_1,\bar{h}_1}(z_1,\bar{z}_1)\cdots\mathcal{O}_{h_n,\bar{h}_n}(z_n,\bar{z}_n)\rangle = \left(\prod_{i=1}^{n}\frac{d\omega_i}{\omega_i}\omega_i^{\Delta_i}\right)\mathcal{A}_n\left(\omega_1\hat{p}(z_1,\bar{z}_1),H_1,n_1;\ldots;\omega_n\hat{p}(z_n,\bar{z}_n),H_n,n_n\right). \tag{35}$$

From this we get our first concrete prediction about the putative dual CFT. For each special of bulk particle with spin s, the boundary theory's spectrum includes a continuous family of conformal primaries with conformal weights

$$\left(h,\bar{h}\right) = \left(\frac{\Delta+s}{2},\frac{\Delta-s}{2}\right). \tag{36}$$

parametrized by the continuum of possible values of Δ. So far it is not clear exactly what range of Δ's should be allowed. One principled statement we can make is that, if we run through this argument in position space for bulk dimension $(d+1,1)$, the conformal primary wavefunctions with weights in the principal series [**?**]

$$\Delta \in \frac{d}{2} + i\mathbb{R} \tag{37}$$

form a complete basis of normalizable solutions to the free wave equation.

3.3 Massive Particles

TODO

4 Dynamics of the Celestial Dual

4.1 OPEs Organized by Symmetry

One of the miracles of 2D CFTs is that, by repeatedly taking operator product expansions (OPEs), one can write the n-point correlator of an arbitrary collection of non-coincident operator insertions in terms of sums and products of lower-point functions. As 2pt and 3pt functions are forced by symmetry up to a finite number of constants, this process almost entirely constrains the dynamics of our theory.

TODO: Write this more clearly, exactly what is left unfixed by OPEs?

On the boundary, computing an OPE means considering a correlator

$$\langle A(z_1,\bar{z}_1)B(z_2,\bar{z}_2)\cdots\rangle \tag{38}$$

where we take $(z_1,\bar{z}_1)$ to be asymptotically close to $(z_2,\bar{z}_2)$, and $\cdots$ hides a finite number of other insertions some finite distance from $(z_2,\bar{z}_2)$. Then inside of our correlator we can replace

$$A(z_1,\bar{z}_1)B(z_2,\bar{z}_2) = \sum_{n,\bar{n}\in\mathbb{Z}}\frac{\{AB\}_{n,\bar{n}}(z_2,\bar{z}_2)}{z_{12}^{n}\bar{z}_{12}^{\bar{n}}} \tag{39}$$

where

$$(z_{ij},\bar{z}_{ij}) := (z_i - z_j, \bar{z}_i - \bar{z}_j) \tag{40}$$

and where the coefficient fields $\{AB\}_{n,\bar{n}}$ are independent of the remaining insertions, are nonsingular at $(z_2,\bar{z}_2)$, and vanish for sufficiently large $n,\bar{n}$. This series converges exactly when $(z_1,\bar{z}_1)$ is closer to $(z_2,\bar{z}_2)$ then any other point in the correlator is. Knowing the OPE between A and B is equivalent to knowing all of the coefficients $\{AB\}_{n,\bar{n}}$.

Now we change notation slightly. We write our operators at $\mathcal{O}_i$, with conformal weights $(h_i, \bar{h}_i)$. By translation covariance we can also expand $(z, \bar{z})$ about the origin rather than $(z_1, \bar{z}_1)$ about $(z_2, \bar{z}_2)$, writing slightly more neatly

$$\mathcal{O}_1(z, \bar{z})\mathcal{O}_2(0,0) = \sum_{n,\bar{n}\in\mathbb{Z}} \frac{\{\mathcal{O}_1\mathcal{O}_2\}_{n,\bar{n}}(0,0)}{z^n \bar{z}^{\bar{n}}}. \tag{41}$$

Covariance under the conformal group strongly constrains the form of the RHS. For one, the LHS and RHS must transform in the same way under conformal transformations. The LHS has conformal weights $(h_1 + h_2, \bar{h}_1 + \bar{h}_2)$

$$\mathcal{O}_1(z, \bar{z})\mathcal{O}_2(0,0) \sim (h_1 + h_2, \bar{h}_1 + \bar{h}_2) \tag{42}$$

while $z^n \bar{z}^{\bar{n}}$ has conformal weights

$$z^n \bar{z}^{\bar{n}} \sim (-n, -\bar{n}) \tag{43}$$

so we can read off the conformal weights of the coefficients as

$$\{\mathcal{O}_1\mathcal{O}_2\}_{n,\bar{n}} \sim (h_1 + h_2 - n, \bar{h}_1 + \bar{h}_2 - \bar{n}). \tag{44}$$

We can make more progress by working in the basis for local operators given by conformal primaries $\mathcal{O}_P$ and their descendants $\partial^n \bar{\partial}^{\bar{n}} \mathcal{O}_P$. Denoting by C_{12}^P the OPE coefficient of the primary $\mathcal{O}_P$ and by $\beta_{12n\bar{n}}^P$ the relative coefficient of the $(n, \bar{n})$th descendant (so by construction $\beta_{1200}^P = 1$), we have

$$\mathcal{O}_1(z, \bar{z})\mathcal{O}_2(0,0) = \sum_P C_{12}^P \sum_{n,\bar{n}} \frac{\beta_{12n\bar{n}}^P}{z^{h_1+h_2-h_P} \bar{z}^{\bar{h}_1+\bar{h}_2-\bar{h}_P}} \partial_n \bar{\partial}_{\bar{n}} \mathcal{O}_P(0,0). \tag{45}$$

This form of the OPE is manifestly covariant under translation, rotation, and scaling, but for totally general coefficients C, β does not yet enforce covariance under special conformal transformations. As it turns out [**?**], further enforcing covariant under SCTs (literally applying $[L_1, \cdot]$ and $[\bar{L}_1, \cdot]$ to both sides and equating the results term-by-term) recursively determinants the β coefficients, leaving only the C's as free data. In fact in a celestial CFT we can do better, since Lorentz symmetry embeds inside of a broader Poincaré group, and this extension on the boundary contrains even the leading coefficients C_{12}^P. More on this later.

4.2 OPEs from Bulk Scattering

We have just seen that boundary global conformal symmetry recursively determines the coefficients of descendants appearing in the OPE between any two primaries, so it suffices to compute just the OPE coefficients C_{12}^P between primaries. We also recall that each continuous family of boundary primaries corresponds to a bulk particle species. Now we make a very nice observation: from our celestial dictionary, we see that taking two insertions $(z_1, \bar{z}_1), (z_1, \bar{z}_2)$ asymptotically close corresponds to taking the corresponding bulk external momenta $p_i = \omega_i \hat{p}(z_i, \bar{z}_i)$ for $i = 1, 2$ asymptotically collinear.

Luckily for us, it is straightforward to see in diagrammatic language why (and when) bulk scattering amplitudes diverge in this limit. If we consider a scattering amplitude with external legs $1, 2$, and some other particle species 3 where $1 + 2 \to 3$ scattering has a nonzero amplitude, then a subset of diagrams contributing to $1 + 2 \to (n - 2)$ scattering will have an internal 3 leg forced by momentum conservation to carry momentum $p_1 + p_2$. In our parametrization, the associated propagator is (modulo convention-dependent constants)

$$\frac{1}{(p_1 + p_2)^2} = \frac{1}{2p_1 \cdot p_2} = \frac{1}{2\omega_1\omega_2 \hat{p}(z_1, \bar{z}_1) \cdot \hat{p}(z_2, \bar{z}_2)} = \frac{-1}{4(z_1 - z_2)(\bar{z}_1 - \bar{z}_2)}. \tag{46}$$

Not only does this explain why boundary coincident singularities make sense from a bulk scattering perspective, it gives us highly quantitative information:

- The primaries P appearing in the boundary $12 \to P$ OPE correspond to bulk particle species with a nonzero scattering amplitude against $1 + 2$, with OPE coefficient scaling with the amplitude for such a process.

- The boundary OPEs are most inverse-linear in z_{12} and $\bar{z}_{12}$.

Conversely, this hints at an analytic subtlety. So far we have been doing all of our complex analysis in a formalism which (at least typographically) treats $z, \bar{z}$ as independent variables. Said more concretely, a complex function which isn't holomorphic may nonetheless admit a local Laurent series involving powers of both z and $\bar{z}$

$$f : \mathbb{C} \to \mathbb{C}, \quad f(z) = \sum_{n,\bar{n}\in\mathbb{Z}} a_{n\bar{n}} z^n \bar{z}^{\bar{n}}. \tag{47}$$

This series expansion then defines a function of two complex variables by taking $z, \bar{z} \in \mathbb{C}$ to be independent. We recover f by working on the "real surface" where $z, \bar{z}$ are conjugates. This is a standard game to play when studying 2D CFTs. On the amplitudes side, however, it raises some concern.

Like OPEs in a 2D CFT, bulk amplitudes are productively studied by continuing off the real surface. Equivalently, if we use spinor helicity notation to factor our null momenta as

$$p_{\alpha\dot{\alpha}} \equiv p_\mu \sigma^\mu_{\alpha\dot{\alpha}} = \lambda_\alpha \tilde{\lambda}_{\dot{\alpha}} \tag{48}$$

it is illuminating to analytically continue to complex momenta by dropping the assumption that $\lambda, \tilde{\lambda}$ be component-wise complex conjugates of one another. But then we have have variously

$$p_1 \cdot p_2 \sim (z_1 - z_2)(\bar{z}_1 - \bar{z}_2) \sim (\lambda_1 \cdot \lambda_2)(\tilde{\lambda}_1 \cdot \tilde{\lambda}_2). \tag{49}$$

So taking $p_1 \cdot p_2 \to 0$ can be accomplished by taking $z_1 - z_2 \sim \lambda_1 \cdot \lambda_2 \to 0$, or $\bar{z}_2 - \bar{z}_2 \sim \tilde{\lambda}_1 \cdot \tilde{\lambda}_2 \to 0$, and these limits may or may not commute. Certainly from the bulk side it's often unclear the correct limit to take, and much of the amplitudes literature contents itself wih taking the holomorphic limit $z_1 - z_2 \sim \lambda_1 \cdot \lambda_2 \to 0$ with the antiholomorphic variables fixed. From this one can recover the opposite limit by symmetry, though recovering the simultaneous limit is more difficult.

In any case, if we are happy to worry just about the holomorphic limit we can use known properties of bulk collinear limits to infer the corresponding boundary OPEs in the same kind of limit. The strategy is as described above: any $1 + 2 \to 3$ channel in the bulk gives rise to a contribution $\mathcal{O}_3$ to the boundary $\mathcal{O}_1\mathcal{O}_2$ OPE with singularity $\sim \frac{1}{z_{12}}$, plus subleading terms corresponding on the boundary to descendants of $\mathcal{O}_3$. The structure of the bulk 3pt vertex then fixes the finer structure of the boundary OPE.

In the simplest example, we can consider tree-level pure-glue Yang-Mills [**?**]. Free of loop-level and non-perturbative corrections, our bulk has only two bulk particle species[1] So tree level $GG \to G$ bulk scattering channel in gives rise to the holomorphic OPE channel

$$\mathcal{O}_{\Delta_1,s_1}(z_1,\bar{z}_1)\mathcal{O}_{\Delta_2,s_2}(z_2,\bar{z}_2) \propto \frac{1}{z_{12}}\mathcal{O}_{\Delta_3,s_3}(z_3,\bar{z}_3) + (\text{descendants}) + (\text{other channels}) + (\text{reg}) \tag{50}$$

where $\mathcal{O}_{\Delta,s}$ for $s = \pm 1$ is the celestial dual to a gluon with helicity s. While a Mellin transform in a vacuum has Δ as a free variable, it turns out that bulk physics fixes Δ_3 in the above OPE. In more detail, a 3pt interaction vertex

$$V_{3\text{pt}} \sim (\partial_{m_1}\Phi_1)\,(\partial_{m_2}\Phi_2)\,(\partial_{m_3}\Phi_3) \quad m_1 + m_2 + m_3 \equiv m \tag{51}$$

has mass dimension $3 + m$[2] This gives rise to a $1 + 2 \to 3$ scattering channel with amplitude scaling like p_3^m. Schematically, we then have a singularity... TODO: come back to this. Need an elementary argument. Anyway we know

$$\Delta_3 = \Delta_1 + \Delta_2 + d_V - 5. \tag{52}$$

[1] Our phrasing here is literal and intentional: under the standard definition of a particle species as a sufficiently small irrep of $SO^+(3,1)$ in our Hilbert space, a positive-helicity gluon cannot be transformed into a negative one. This point is often glossed over.

[2] A bosonic field with a standard kinetic term in D dimensions has mass dimension $\frac{D-2}{2}$, which in $3 + 1$ reduces to 1.